

KEHINDE (KENNY) SAMSON OGUNLOWO

Principal AI Platform and Multi-Cloud Security Architect

Generative and Agentic AI | LLMOps | AI Security and Governance | AWS, Azure, GCP

Houston, TX | 347.757.8099 | kogunlowo@gmail.com | [LinkedIn](#) | [GitHub](#) | [Book a call](#) | **Active Secret Clearance**

PROFESSIONAL SUMMARY

Principal AI Platform Architect and Multi-Cloud Security Architect with 11+ years across healthcare, defense, financial services, and energy. Designs and builds secure, cloud native AI platforms that let organizations deploy, govern, and operate Large Language Model (LLM) and agentic AI systems in production: multi-agent orchestration, Retrieval Augmented Generation (RAG) pipelines, model evaluation, prompt guardrails, Infrastructure as Code (IaC), CI/CD security scanning, Policy as Code, and workload identity. Builds in Python end to end: FastAPI services, async agent runtimes, retrieval and data pipelines orchestrated with Apache Airflow, Model Context Protocol (MCP) tool servers, and evaluation harnesses, with governed analytics across Snowflake, BigQuery, and Azure Synapse.

Principal level architecture with staff engineer execution: owns the design, writes the Python and the Terraform, wires the CI/CD, builds the orchestration, and stays hands on until the system is live and measured. Communicates architecture in business terms: systems, guarantees, constraints, and measurable outcomes. Covers the AI engineering lifecycle across AWS (Bedrock, SageMaker), Azure (Azure OpenAI, AI Foundry), and GCP (Vertex AI), with LangChain and LangGraph orchestration behind LLM gateways and eval gated delivery, on programs aligned to FedRAMP High, CMMC Level 2, HIPAA, HITRUST, SOC 2, PCI-DSS, and NIST 800-53 with Zero Trust architecture.

CORE EXPERTISE

AI Platform Engineering and Agentic AI

Enterprise AI platforms, autonomous agents, Retrieval Augmented Generation (RAG), LLMOps and MLOps, AI evaluation and observability, multi-agent orchestration, Model Context Protocol (MCP) and Agent2Agent (A2A), LangChain, LangGraph, AutoGen, CrewAI, Python, FastAPI, async agent runtimes, REST and OpenAPI contracts.

Built the enterprise RAG and multi-provider model routing layer grounding tax automation agents across 500K+ transactions per month.

Cloud, Data, and DevSecOps Architecture

AWS, Azure, GCP, AWS GovCloud, Azure Government, Kubernetes, Docker, Terraform, Infrastructure as Code (IaC), CI/CD, GitHub Actions, GitOps, platform engineering and developer experience, data pipeline orchestration with Apache Airflow, analytics on Snowflake, BigQuery, and Azure Synapse, streaming with Apache Kafka.

Cut CI/CD security vulnerabilities 75% quarter over quarter across 50+ microservices.

AI Security and Cybersecurity Architecture

LLM security and prompt injection defense, Zero Trust, IAM, workload identity federation, Policy as Code, SIEM and SOAR, detection engineering, CNAPP and CSPM (Wiz), endpoint and cloud workload protection (CrowdStrike Falcon), cloud security automation and auto remediation, confidential computing.

Delivered environments aligned to FedRAMP High, StateRAMP, HIPAA, NIST 800-53, NIST AI RMF, and CMMC Level 2.

AI AND CLOUD SECURITY DOMAIN COVERAGE

Hands on coverage across the AI and cloud security surface, applied on regulated production platforms and in published engineering references.

- **Model and Application Layer:** AI Security, LLM Security, Agentic AI Security, MCP Security, Prompt Security, RAG Security, Vector Database Security, Model Security, API Security.
- **Platform and Infrastructure:** Identity Security, Zero Trust Architecture, Cloud Security across AWS, Azure, and GCP, Kubernetes Security, Data Security, Supply Chain Security, DevSecOps, LLMOps Security.
- **Governance and Assurance:** AI Governance, Responsible AI, AI Risk Management, AI Compliance, Continuous Compliance, Privacy Engineering, AI Audit and Forensics.
- **Offensive and Operations:** AI Threat Modeling, AI Red Teaming, AI Observability, Incident Response, Security Automation.
- **Engineering Controls:** Policy as Code, Infrastructure as Code Security, Secure CI/CD.

SELECTED IMPACT

- **40% reduction in manual clinical review** using LLM document workflows and RAG on a platform serving 15M+ members and 2M+ claims per month (Cigna).
- **9x faster threat detection:** mean time to detect cut from 45 minutes to under 5 through SIEM modernization and custom detection engineering (Cigna).
- **70% faster provisioning** across 12 production environments with Terraform first landing zones and self service golden paths (Ceretax).
- **25% multi-cloud cost reduction** through landing zone standardization and multi-provider model routing tuned to cost, latency, and quality (Ceretax, Lockheed Martin).
- **Zero high severity audit findings** on regulated platforms clearing FedRAMP High and StateRAMP (Ceretax, Lockheed Martin).

PROFESSIONAL EXPERIENCE

Principal Multi-Cloud Security Architect and AI Engineer | Ceretax, Houston, TX

January 2024 to Present

Tax compliance fintech | 25 engineer cross functional platform organization

User benefit: tax and finance teams stopped hand researching determinations. Every answer came back with its source attached, so a reviewer could accept or challenge it in seconds instead of rebuilding the research, and the same citation dropped straight into the audit file when the assessor asked how the number was reached.

- Built the enterprise multi-cloud (AWS, Azure, GCP) Infrastructure as Code platform in Terraform, Bicep, ARM, and Cloud Deployment Manager with repeatable landing zones, embedding LLM driven tax automation on Azure OpenAI GPT-4 and Python multi-agent frameworks (LangChain, LangGraph, AutoGen, CrewAI) on GKE Autopilot, cutting provisioning time 70% across 12 production environments.
- Stood up the enterprise RAG layer in Python on PostgreSQL full text search (tsvector and GIN) fused with pgvector, with reciprocal rank fusion, cross encoder reranking, parity gates, and a citation on every answer, with Apache Airflow orchestrating corpus ingestion and index refresh, grounding tax automation agents at audit grade traceability across 500K+ transactions per month.
- Built the multi-provider model routing layer as a FastAPI LLM gateway exposing a versioned REST API across Azure OpenAI, AWS Bedrock, and Vertex AI, with Pydantic request contracts, token budgets, and failover so each task class runs on the lowest cost model that clears its quality bar, and exposed internal systems to agents through Model Context Protocol (MCP) tool servers with scoped tool invocation controls.
- Engineered the Zero Trust GKE security architecture hands on with Workload Identity, Binary Authorization, Pod Security Standards, Istio mTLS, Identity Aware Proxy, VPC Service Controls, and Cloud Armor, clearing FedRAMP High and StateRAMP audits with zero high severity findings at a 99.99% uptime SLA.
- Wired secure CI/CD in Cloud Build, GitHub Actions, and Azure DevOps with Policy as Code (OPA Gatekeeper, Config Validator), SAST and DAST gates via GitHub Advanced Security, and Secret Manager integration, then closed the loop by routing Wiz CNAPP findings and CrowdStrike Falcon cloud workload protection alerts into event driven auto remediation, driving vulnerabilities down 75% quarter over quarter across 50+ microservices.
- Established FIPS 140-2 validated encryption on Cloud KMS with HSM backing, managing 200+ encryption keys for sensitive financial data under federal regulatory requirements.
- Instrumented end to end AI observability and LLM Ops pipelines in Python covering model lifecycle, drift detection, evaluation harnesses, and distributed tracing, warehousing cost and token analytics in Snowflake so model deploys travel the same disciplined secure path as the rest of the platform.
- Implemented the AI governance baseline aligned to NIST AI RMF, coding Python inference REST API security controls, drift monitoring thresholds, and prompt injection guardrails on production LLM workloads, and set the quarterly cloud and AI roadmap with the VP of Engineering, CISO, and product leadership.

Staff Multi-Cloud Security Architect and Healthcare AI Engineer | Cigna, Remote

February 2023 to December 2023

22 engineer healthcare platform team | 15M+ member population

User benefit: nurse reviewers opened a case with the policy language, the clinical evidence, and the citations already assembled instead of searching four source systems for them. The clinician still made every medical necessity call, but spent the time on judgment rather than retrieval, cutting manual review effort 40% across 2M+ claims per month and returning decisions faster to the member and provider waiting on them.

- Built and deployed Python LLM automation with LangChain and RAG pipelines on Azure OpenAI GPT-4 for claims processing, prior authorization, and medical necessity determination, cutting manual review time 40% across 2M+ claims per month.
- Engineered the Python hybrid retrieval layer pairing PostgreSQL full text search with vector search and semantic reranking so clinical answers stayed grounded and auditable, with parity gates that caught retrieval regressions before they reached a clinician.
- Built a real time clinical analytics platform on Azure Synapse, BigQuery, and Snowflake, with Apache Airflow orchestrating ingestion and transformation pipelines processing 10TB+ of daily healthcare data mapped to SNOMED CT, LOINC, and ICD-10 terminologies for population health insights.
- Migrated FHIR R4 services to Azure Health Data Services and Google Cloud Healthcare API, building secure REST API gateways interoperating with Epic, Cerner, and Meditech under HIPAA and HITRUST control sets, unlocking 8 net new clinical partner integrations.
- Cut mean time to detect for PHI exfiltration, ransomware, and insider threats from 45 minutes to under 5, a 9x improvement, by deploying Chronicle SIEM and SOAR alongside Microsoft Sentinel, Splunk, and CrowdStrike Falcon EDR telemetry and authoring the custom detections in Python and KQL, and ran the incident response playbooks those detections fed.
- Engineered the FedRAMP High migration to Google Cloud with Assured Workloads, FIPS 140-2 PHI encryption, and continuous monitoring via Security Command Center, compressing audit preparation from 6 weeks to 5 days.
- Implemented responsible AI guardrails in Python for clinical LLM use cases: PHI redaction, hallucination detection, citation validation, and human in the loop escalation gates on medical necessity determinations.

Lead Multi-Cloud Security Architect and AI/ML Platform Engineer | Lockheed Martin, Remote and Hybrid

January 2021 to January 2023

18 engineer cleared multi-cloud platform team | 30+ programs

User benefit: maintainers and command staff saw which of 1,000+ connected assets needed attention before it failed, with multiple sensor feeds fused into one picture instead of several consoles to reconcile by hand. The data never left the accreditation boundary to get there, so the people closest to the mission got the answer without a classification review standing in the way.

- Delivered the FedRAMP High and CMMC Level 2 cloud migration for mission critical defense workloads across Azure Government, AWS GovCloud, and Google Cloud with multi-region failover, reducing infrastructure cost 25% and briefing program managers and government sponsors on compliance posture and mission SLAs.
- Engineered Zero Trust segmentation, hardware rooted attestation, and continuous monitoring for cleared workloads held to standards that survive ATO audit, partnering with government sponsors on authorization posture across 30+ programs.
- Architected Python based AI and ML systems for sensor data fusion, logistics optimization, and predictive maintenance, processing real time IoT telemetry from 1,000+ connected assets through Azure Synapse and BigQuery with embedded Power BI dashboards for command staff decision support.
- Engineered confidential computing environments using Intel SGX, AMD SEV, Confidential GKE, and AWS Nitro Enclaves with hardware rooted attestation to protect classified data during processing.
- Built the secure CI/CD and GitOps delivery path, with Python automation, that shipped models and REST services into IL4 environments with Policy as Code gates, image scanning, and software supply chain controls.
- Built enterprise key management on Azure Dedicated HSM, Cloud HSM, and AWS CloudHSM, implementing BYOK and HYOK strategies for cryptographic operations across the program portfolio.

Senior Multi-Cloud AI/ML Platform Engineer, Healthcare Solutions | Catalyte and NantHealth, Remote *January 2020 to December 2020*

- Built healthcare AI and ML pipelines in Python on Amazon SageMaker, Comprehend Medical, and Azure Cognitive Services for clinical decision support, EMR and EHR NLP, automated patient triage, and clinical documentation improvement, increasing triage accuracy 30%.
- Operationalized clinical NLP models in Python for documentation improvement and coding support, wiring SNOMED CT, LOINC, ICD-10, and RxNorm terminologies into the inference path for standardized, structured output.
- Engineered FHIR compliant ETL pipelines in Python on Azure Data Factory, AWS Glue (PySpark), and Apache Airflow ingesting HL7, DICOM, and CDA data for HIPAA compliant real time health information exchange, and load tested telehealth platforms on AKS with KEDA autoscaling to sustain 50K+ concurrent users at a 99.9% SLA.

Senior Cloud Security Architect and AI Solutions Engineer | Mammoth Energy Services, Oklahoma City, OK *October 2019 to January 2020*

- Built Python event driven automation with Azure Logic Apps, Functions, AWS Lambda, and SageMaker for predictive maintenance and equipment scheduling, cutting operational costs 30%.
- Stood up the observability backbone across geographically distributed field operations on the ELK stack with Prometheus and Grafana, processing 5TB+ of daily operational data into a single real time operational view for remote sites.

Senior Cloud Security Architect | BP Refinery, Texas City, TX *August 2018 to September 2019*

- Automated enterprise vulnerability management across 5,000+ host assets by wiring Qualys and Rapid7 into CI/CD, and embedded Python SageMaker anomaly detection into pipelines for predictive security analysis.
- Drove AWS, Docker, Kubernetes, and Ansible infrastructure automation that cut deployment time 50%, hardening container and cluster build pipelines with image scanning and codifying infrastructure for repeatable, audited deployments.

Senior Cloud Infrastructure Architect and AI/ML Engineer | Patterson UTI, Houston, TX *March 2017 to July 2018*

- Architected drilling optimization with Azure Machine Learning and predictive analytics, reducing non productive time and operational cost 20%.
- Stood up Jenkins CI/CD that improved release cycles 30%, and automated multi-cluster Kubernetes (Kafka, Elasticsearch) lifecycle with Terraform, Ansible, PowerShell, and Puppet for repeatable operations.
- Instrumented real time telemetry ingestion and Python feature pipelines feeding the predictive models, and built the dashboards that surfaced non productive time drivers to drilling operations.

Cloud Automation and Infrastructure Engineer | Securitas USA, Houston, TX *April 2015 to March 2017*

- Streamlined enterprise CI/CD with Ansible Tower, Jenkins, and Maven, cutting deployment time 40%, and ran AWS infrastructure (EC2 fleet, ELB, Auto Scaling) at 99.9% availability with containerized microservices serving 10,000+ users.
- Implemented Azure Backup and Site Recovery and auto scaling across the AWS estate to meet enterprise RTO and RPO targets while optimizing infrastructure cost 25%, standardizing release automation so deployments were repeatable and rollback safe.

INDEPENDENT PROJECTS AND PUBLISHED REFERENCES

Personal work built and published outside of employment. These are public repositories and engineering reference documents, not employer deliverables or production deployments.

Citadel SaaS Factory | Open source multi-agent framework | github.com/Citadel-Cloud-Management/citadel-saas-factory

- Author and maintainer of an open source Python agent framework: 265 agent definitions across 15 SaaS domains, with multi-LLM routing across 12 model providers (paid, free tier, and local) tuned to cost, latency, and capability per task class.

- Designed the canonical agent factory in Python: BaseAgent abstraction with Pydantic typed contracts and a defined execution lifecycle, FastAPI service surface with a versioned REST API and MCP tool servers, multi-provider routing with failover and token budgets, a RAG layer that ships citations on every answer, multi-agent orchestration with shared memory and consensus, guardrails (PII detection, prompt injection defense, output validation), and packaged deployment patterns covering canary, blue green, rollback, and multi-region serving.

Enterprise AI Agent Platform | Cloud agnostic reference architecture | AWS, Azure, GCP

- Authored a unified reference architecture for a cloud agnostic enterprise AI agent platform: Terraform contract layer modules with parallel AWS, Azure, and GCP implementations behind identical interfaces, a shared RAG spine (hybrid lexical and vector retrieval, RRF fusion, cross encoder reranking, ACL filtered results), a single FastAPI LLM gateway chokepoint exposing one REST API for routing, failover, semantic caching, and per tenant cost attribution, and first class AppOps, NetOps, SecOps, DevOps, and DataOps planes enforced through CODEOWNERS.
- Catalogued 393 agent definitions across 24 business functions and 7 domains, each a declarative registry entry binding corpora, retrieval profile, MCP tool grants, guardrails, sha256 pinned prompts, and a T0 to T3 risk tier gated by evaluation suites, so a new business agent is configuration plus evals rather than new code paths.

Agentic RAG Platform | Engineering reference and operating manual | Local MVP and AWS Bedrock

- Documented a hexagonal RAG architecture covering eight retrieval patterns, retrieval latency budgets, an eval framework (golden set, LLM judges, PR gate), immutable prompt versioning with rollback, end to end tenant isolation, disaster recovery with RTO and RPO drills, a cost model, guardrails, NIST 800-53 Rev 5 control mapping, and 20 architecture decision records.
- Implemented the hybrid retrieval path end to end in Python across a local MVP and AWS Bedrock, served behind a FastAPI REST API: query intent, lexical index, vector semantic embeddings, fuse and rerank scoring, and grounded cited answers, with evaluation and parity gates that catch retrieval regressions before release.

Multi-Cloud AI Security Field Manual | 20 domain AI security reference | AWS, Azure, GCP

- Authored control coverage across 20 AI security domains including prompt injection defense, RAG security, agentic AI security, AI supply chain and runtime security, MCP and tool invocation security, tenant isolation, AI red teaming, AI threat detection, and AI DLP, with each control mapped to its AWS, Azure, and GCP native service, a cross role RACI, and a production hardening runbook.
- Documented the AI threat modeling, AI red teaming, and AI audit and forensics playbooks in the same reference: adversarial test suites for prompt injection and jailbreak, model and vector store abuse cases, evidence retention for agent and tool invocation traces, and continuous compliance evidence mapped to NIST AI RMF and NIST 800-53 Rev 5.

TRIPLE-CLOUD REFERENCE ARCHITECTURE (DESIGN REFERENCE)

One tiered web and REST API topology (Python FastAPI services) expressed behind identical Terraform contract layer interfaces on all three clouds: DNS, CDN, WAF, L7 load balancer, managed Kubernetes, relational, cache, streaming, object and shared file storage, with telemetry, tracing, workload identity, and secrets on the platform plane.

- **AWS:** Route 53 > CloudFront > AWS WAF > Application Load Balancer > Amazon EKS > Amazon RDS, ElastiCache for Redis, Amazon MSK, Amazon S3, Amazon EFS, with CloudWatch, X-Ray, IAM and IRSA, Secrets Manager, and KMS.
- **Azure:** Azure DNS > Azure Front Door > Azure WAF > Application Gateway > Azure Kubernetes Service > Azure Database for PostgreSQL Flexible Server, Azure Managed Redis, Event Hubs, Blob Storage, Azure Files, with Azure Monitor and Log Analytics, Application Insights, Entra ID Workload Identity, and Key Vault.
- **GCP:** Cloud DNS > Cloud CDN > Cloud Armor > Global External Application Load Balancer > GKE Autopilot > Cloud SQL and AlloyDB, Memorystore for Redis, Managed Service for Apache Kafka and Pub/Sub, Cloud Storage, Filestore, with Cloud Monitoring and Logging, Cloud Trace, Workload Identity Federation, and Secret Manager.
- **Cross-cutting:** Istio mTLS and Pod Security Standards, Policy as Code admission gates, private egress (PrivateLink, Private Endpoints, Private Service Connect), multi-region failover, and per tenant cost attribution.

TECHNICAL SKILLS

- **Languages, APIs, and Frameworks:** Python (primary: FastAPI, Pydantic, asyncio, LangChain, pandas, boto3, pytest), REST and OpenAPI design, Model Context Protocol (MCP) servers and clients, gRPC, GraphQL consumption, SQL, Bash, PowerShell, HCL and Terraform, YAML, KQL
- **AI and LLM Platforms:** Azure OpenAI (GPT-4), Azure AI Foundry, AWS Bedrock (Claude, Llama, Titan), Google Vertex AI (Gemini, PaLM 2), Amazon SageMaker, Azure Machine Learning, Hugging Face, TensorFlow, PyTorch, Amazon Comprehend Medical, foundation models, embeddings, fine tuning
- **Agentic AI and Orchestration:** Python, FastAPI, Pydantic, asyncio, LangChain, LangGraph, LangSmith, AutoGen, CrewAI, Amazon Bedrock Agents, Azure AI Agent Service, Vertex AI Agent Builder, Model Context Protocol (MCP) servers and clients, Agent2Agent (A2A), multi-provider model routing, MCP tool catalogs, agent failure mode handling, LLM tracing and evaluation
- **Agent Design and Runtime:** ReAct, chain of thought, plan execute reflect, self critique and verification, task decomposition and planning, working and long term memory, retry and replanning with backoff, human in the loop approval gates, parallel execution, timeouts and limits, rollback and recovery
- **RAG and Vector Search:** Python retrieval services, PostgreSQL full text search (tsvector and GIN), pgvector, Azure AI Search, Amazon Bedrock Knowledge Bases, Vertex AI Search, OpenSearch Serverless, hybrid retrieval with reciprocal rank fusion and cross encoder

rerank, Apache Airflow orchestrated corpus ingestion and refresh, golden set eval and parity gates, hallucination detection, citation validation, prompt versioning and rollback

- **AI Security and LLM Security:** Prompt security and prompt injection defense, jailbreak and adversarial input testing, RAG security (corpus ACLs, ingestion trust boundaries, citation validation), vector database security (tenant scoped indexes, embedding leakage and inversion controls), model security (artifact signing, provenance, serialization safety), agentic AI security (tool scoping, autonomy limits, action approval gates), MCP security (server allowlists, tool invocation controls, credential brokering, audit trails), API security for inference endpoints (authn and authz, rate and token limits, schema validation), AI red teaming, AI threat modeling, AI DLP and PHI redaction
- **AI Evaluation and Quality:** Task success rate, accuracy and correctness, hallucination rate, groundedness, reliability and robustness, latency and throughput, cost efficiency, human evaluation, A/B testing, continuous evaluation, evals as CI gates
- **LLMOps and AI Governance:** LLMOps and MLOps pipelines, model lifecycle management, AI observability, cost and token analytics, replay and debugging, inference API security, model drift detection, NIST AI RMF, ISO/IEC 42001 alignment, NIST 800-53 Rev 5 control mapping for AI workloads, AI risk management and risk tiering, responsible AI, AI compliance and continuous compliance evidence, AI audit and forensics (agent and tool call trace retention, replay for investigation), prompt injection guardrails, MCP tool invocation controls, sandboxing and isolation, data leakage prevention, human in the loop gates, semantic caching, LLM gateways
- **Cloud Networking and Edge:** Route 53, Cloud DNS, Azure DNS and Traffic Manager, CloudFront, Cloud CDN, Azure Front Door, AWS WAF and Shield, Azure WAF, Cloud Armor, Application Load Balancer, Application Gateway, Global External Application Load Balancer, PrivateLink, Private Endpoints, Private Service Connect, VPC design and segmentation
- **Cloud Security and Zero Trust:** Wiz (CNAPP, CSPM, CIEM, DSPM, container and IaC scanning), CrowdStrike Falcon (EDR and XDR, Cloud Security Posture Management, Cloud Workload Protection, Identity Protection, threat intelligence), BeyondCorp Enterprise, Identity Aware Proxy, VPC Service Controls, Security Command Center, Cloud Armor, Microsoft Defender XDR, Microsoft Sentinel, Chronicle SIEM and SOAR, Splunk, GuardDuty, Security Hub, Macie, AWS Config, Prisma Cloud, confidential computing (Intel SGX, AMD SEV, AWS Nitro Enclaves, Confidential GKE)
- **Cloud Security Automation:** Wiz and CrowdStrike Falcon API and webhook integrations, event driven auto remediation (EventBridge, Cloud Functions, Azure Logic Apps), SOAR playbooks and case automation, guardrails as code with OPA and Terraform Sentinel, drift detection and automated rollback, misconfiguration and vulnerability auto ticketing, continuous compliance evidence collection, Python and Terraform automation across AWS, Azure, and GCP
- **Kubernetes and Container Security:** Pod Security Standards, admission control with OPA Gatekeeper, Binary Authorization and image signing, workload identity, network policy and Istio mTLS, runtime and image scanning, secrets management, node and cluster hardening on GKE Autopilot, AKS, EKS, and OpenShift
- **Identity and Cryptography:** Microsoft Entra ID, AWS IAM, Google Cloud IAM, SAML 2.0, OAuth 2.0, OpenID Connect, CyberArk PAM, Azure AD PIM, AWS KMS and CloudHSM, Azure Key Vault and Dedicated HSM, Cloud KMS and HSM, BYOK and HYOK, FIPS 140-2, TLS 1.3, mTLS
- **DevSecOps and IaC Security:** Terraform, Bicep, ARM, CloudFormation, Cloud Deployment Manager, Helm, GitHub Actions, Azure DevOps, Cloud Build, Jenkins, Ansible, GitOps (ArgoCD), Policy as Code (OPA Gatekeeper, Terraform Sentinel, Config Validator), IaC scanning and drift detection, Binary Authorization, Prisma Cloud, GitGuardian, SAST and DAST, secure CI/CD gates, software supply chain security (SBOM, artifact signing, dependency and image scanning), security automation
- **Data Security and Privacy Engineering:** Data classification and tenant isolation, encryption at rest and in transit, key custody and rotation, PHI and PII redaction, data minimization and retention, DLP, access review and least privilege, GDPR and HIPAA privacy controls
- **Detection, Response, and Forensics:** Detection engineering in Python and KQL, SIEM and SOAR automation, CrowdStrike Falcon detection and response workflows, Wiz finding triage and prioritization, MITRE ATT&CK mapping, STIX threat intelligence, incident response playbooks and tabletop exercises, forensic evidence retention, continuous monitoring (ConMon)
- **Containers and Observability:** GKE Autopilot, AKS, EKS, OpenShift, Docker, Kubernetes, K3s, Istio service mesh, KEDA, Prometheus, Grafana, OpenTelemetry, Azure Monitor, Application Insights, Cloud Trace, AWS X-Ray, CloudWatch, ELK stack, SLO and SLI dashboards, distributed tracing, alerting and anomaly detection, GPU utilization, token economics
- **Data Engineering, Storage, and Streaming:** Snowflake, Apache Airflow, ETL and ELT pipeline design, Apache Kafka, Amazon MSK, Azure Event Hubs, Google Managed Service for Apache Kafka, Pub/Sub, EventBridge, Amazon RDS, Cloud SQL and AlloyDB, Azure Database for PostgreSQL, ElastiCache, Azure Managed Redis, Memorystore, Amazon S3, Blob Storage, Cloud Storage, Amazon EFS, Azure Files, Filestore, AWS Glue, Azure Data Factory, Dataflow, BigQuery, Amazon Security Lake, Azure Synapse, Azure Databricks, OpenSearch, Elasticsearch
- **Healthcare and Domain AI:** FHIR R4, HL7, DICOM, Azure Health Data Services, Google Cloud Healthcare API, Epic, Cerner, Meditech, SNOMED CT, LOINC, ICD-10, RxNorm, CPT, CDS Hooks, claims and prior authorization automation
- **Enterprise AI Integrations:** Slack, Microsoft 365, Google Workspace, Salesforce, SAP, ServiceNow, Jira, GitHub, Notion, custom APIs, Amazon Q Business and Developer, Lex, QuickSight Q, HealthLake, Microsoft 365 Copilot, Fabric and Power BI Copilot, Gemini Code Assist, Dialogflow CX, Looker AI, Document AI

CERTIFICATIONS

- **Cloud and Architecture:** Google Cloud Professional Security Engineer, Google Cloud Professional DevOps Engineer, AWS Certified Solutions Architect Professional, AWS Certified DevOps Engineer Professional, Microsoft Certified Azure Solutions Architect Expert, Azure Security Engineer Associate
- **Security and Compliance:** Certified Ethical Hacker (CEH) v11, CompTIA CySA+, CompTIA Cloud+, Palo Alto Networks Certified Network Security Engineer
- **Identity and Automation:** Microsoft SC-300 Identity and Access Administrator Associate, HashiCorp Certified Terraform Associate

COMPLIANCE AND REGULATORY FRAMEWORKS

FedRAMP High, CMMC Level 2, StateRAMP, NIST 800-53 Rev 5, NIST 800-171, NIST AI RMF, DoD SRG IL4, ITAR, HIPAA, HITRUST, SOC 2 Type II, PCI-DSS Level 1, ISO 27001, GDPR, FIPS 140-2, Continuous Monitoring (ConMon), Google Assured Workloads.

EDUCATION

- Bachelor of Applied Mathematics | Tai Solarin University of Education, Nigeria
- Associate Degree in Computer Science | Institute for Industrial Technology, Lagos, Nigeria